

Grounding Documents for Long-Running Tabletop Campaigns: A Practitioner’s Survey of Twelve Approaches, from LLM Extraction to Authored Structure

Kostadis Roussos
Independent researcher
kostadis@gmail.com

October 2026

Abstract

A game master who uses language models to prepare sessions of a long-running tabletop role-playing campaign needs *grounding documents*: a current world state, campaign state, party record, planning notes and per-character dossiers, distilled from dozens of recorded sessions and fed to the models as context. Producing these documents is easy to start and hard to get right, because the details that matter are exactly the ones generative models get wrong: who did what, in what order, which name denotes which entity, and whether something happened at all. Over seven months (March–October 2026) and six campaigns of up to 70 chapters, one practitioner built, measured and retired twelve distinct approaches to the problem. We reconstruct them from the git history, design documents and experiment write-ups of the repositories involved, and organise them by two axes: the *substrate* a pipeline reads (narrated chapter prose, atomic extracted facts, retrieval indexes, or the structured session summaries the GM already reviews) and *who makes the precision decisions* (ordering, attribution, identity, scope). Approaches that let a model make those decisions failed in ways that were silent and that compounded: a consolidated extraction shrank a planning document from 513 to 109 lines; an atomic-fact ensemble reported a slain antagonist as alive, and 98% of its 15,465 facts were seen by only one sample, so agreement could not serve as a filter; a knowledge-graph library stored 1.18 characters of generated text per character of source and answered “what was the most recent session” with session 8 of 11. The approach that the GM finally promoted parses the scene order, entity sections and memorable moments already authored in reviewed session summaries, with no model call (67 summaries, 409 scenes, 868 dossiers, byte-identical across runs), and uses a model only to render the final prose; it also flagged 51 claims in the GM’s own tracking files as unsupported by the record. We distil the history into ten cross-cutting findings and a set of design rules for grounding long-horizon generative work.

1 Introduction

Tabletop role-playing campaigns are long. A weekly game runs for years, and the record of what the players have done—who they met, what they promised, which villain is dead and which only appears to be—is the substrate from which the game master (GM) prepares every subsequent session. When the GM asks a language model to help with preparation (draft an encounter, voice a returning character, check a plot thread), the model needs that record as context. Feeding it hundreds of thousands of words of transcript is impractical and, past a point, counter-productive [11]. The practical answer is a small set of *grounding documents*: compact, current, authoritative summaries of the world, the campaign, the party and the plans, together with per-entity dossiers.

This paper surveys how one practitioner tried to build those documents. The work was done in an open-source toolkit for campaign management (CampaignGenerator [1]), a private repository of campaign data for six campaigns, a public repository of local-model experiments on two NVIDIA DGX Spark machines [2], and a number of local-only scratch directories. Several approaches are still live, several survive only in history, and some never reached version control. The aim is to make that history usable: to describe each approach precisely enough that it can be compared with the others, to record the measured reason each one was kept or dropped, and to extract what generalises.

The problem looks like summarisation, and the first approaches treated it that way. It is better understood as *record keeping under precision constraints*. A grounding document is consumed by another model that cannot tell a confident error from a fact. Errors in four kinds of decision are both likely and expensive: *ordering* (what happened before what), *attribution* (who did it, to whom), *identity* (which names denote the same entity) and *scope* (what belongs in this document, and what actually happened as opposed to what was planned). The toolkit’s governing rule, stated before any of the work surveyed here, is that a model may draft but these *precision decisions* require a human checkpoint before they feed another model. Much of the history below is the gradual discovery of where that rule had been violated without anyone noticing, and of how to remove the model from those decisions rather than merely review its output.

Contributions.

- A catalogue of twelve approaches to producing grounding documents for long-running campaigns, reconstructed from primary sources, with status, lineage and measured outcomes (Sections 4–5).
- A two-axis taxonomy (substrate $\times$ locus of precision decisions) under which the successes and failures separate cleanly (Section 4).
- Ten cross-cutting findings, each with a measurement, that apply beyond tabletop games to any long-horizon generative pipeline that must keep a record (Section 6).

2 Background

2.1 The campaigns and their record

Six campaigns contributed evidence. They are adaptations of published adventures: *Out of the Abyss* (70 chapters at the time of writing), a hybrid of *Dragon of Icespire Peak* and *Lost Mine of Phandelver* (53 sessions), *Storm King’s Thunder*, *Temple of Elemental Evil*, an organised-play adventure set in Hillsfar, and a second *Lost Mine* table (12 sessions). Each session is recorded and transcribed. A post-session pipeline produces, in order: a reviewed transcript (spelling and speaker attribution corrected by the GM); a structured *session summary* with fixed sections (**Summary**, **Memorable Moments**, an ordered **Scenes** list, and entity sections for NPCs, locations, items and spells); scene extractions; and finally a narrated prose chapter. The narrated chapters are concatenated into a campaign *bible* of several megabytes.

2.2 The grounding documents

Four documents form the core: *world state* (the setting as it now stands), *campaign state* (where the story is, open threads, recent events), *party* (characters, resources, arcs) and *planning* (NPC dossiers, factions, threat tracking). They are read by the toolkit’s preparation agents and by interactive

assistants. A per-campaign *entity registry* records canonical names, types and approved aliases for about a thousand entities per campaign; it is the toolkit’s authority on identity and spelling and, as Section 6 argues, it is deliberately a names layer, not a state layer.

2.3 Related work

Retrieval-augmented generation [6] and graph-structured retrieval [7] are the standard alternatives to static context documents; tiered memory systems for agents [8, 9] maintain summaries and reflections that are themselves model-written. Book-length summarisation [10] measures the coherence errors that accumulate when a model compresses a long narrative incrementally. Work on language models in *Dungeons & Dragons* has used actual-play logs and structured game state as training and evaluation data [12, 13] and has built assistants for the dungeon master at the table [14]. None of these addresses the maintenance problem surveyed here: keeping a reviewed, provenance-bearing record of a campaign current over months, under a rule that a human, not a model, settles ordering, attribution and identity. The hallucination literature [15] catalogues the failure modes; this survey records which pipeline shapes expose a practitioner to them and which do not.

3 Survey method

The sources are (i) the full git history of the toolkit repository, including all 28 local and remote branches, deleted files and merged pull requests; (ii) the campaign-data repository (two clones, about 130 branches); (iii) the experiment repository and eleven further repositories or forks, of which only the practitioner’s own commits were considered; (iv) local directories outside version control; and (v) the design documents and experiment write-ups committed alongside the code, which record measurements and kill decisions at the time they were made. Three independent passes catalogued the repositories, and their reports were reconciled against the primary documents; every number in this paper is quoted from a design document, experiment write-up or commit message.

An approach is counted as *distinct* if it changes the substrate, the locus of a precision decision, or the output form; a parameter change within an approach is not counted. Status is one of: *live* (on the main branch and used for at least one campaign’s current documents); *dormant* (on main, not used); *history* (deleted, recoverable from git); *proposal* (designed and measured, not built); *local* (exists only outside version control).

4 A taxonomy

Two axes separate the approaches.

Substrate. What the pipeline reads: (P) narrated *prose*, the bible chapters; (F) atomic *facts* extracted from prose by a model; (R) a *retrieval* index over some corpus, queried at preparation time; (S) the *structured* session summaries that the GM already reviews as part of the post-session pipeline. The narrated prose is downstream of the summaries: a pipeline that reads prose to recover structure is undoing a rendering step.

Locus of precision decisions. Who decides ordering, attribution, identity and scope: a model (M); deterministic code over declared structure (D); or the GM through a recorded ruling (G). Most approaches mix loci; Table 1 gives the dominant one for the step that determines what the document asserts.

Table 1: The twelve approaches. Substrate: P prose, F atomic facts, R retrieval index, S structured summaries. Locus: who makes the ordering/attribution/identity/scope decisions—M model, D deterministic code, G GM ruling.

#	Approach	Subst.	Locus	Active	Status
1	Per-document extract-and-synthesise scripts	P	M	Mar–Aug	live (3 campaigns), dormant elsewhere
2	Shared chapter-extract consolidation	P	M	Apr	history (killed)
3	Hand curation and dossier merge	P	G	Apr–	live as stopgap
4	Local atomic-fact ensemble	F	M	May–Sep	live, no longer feeds OOTA
5	Ensemble repairs: recency, offsets, coverage	F	D/G	Jul	live (within 4)
6	Two-stage finder/claim extraction	F	D	Sep	proposal
7	Entity-level typing and ruling replay	F	D/G	Oct	proposal, replayed
8	Model work over ensemble dossiers	F	M	Jul–Aug	local
9	State stores and per-section projections	F	D/G	Jul–Aug	dormant
10	Knowledge-graph and semantic retrieval	R	M	Jul–Sep	dormant (graph); live (search)
11	Event-ordinal graph over authored scenes	S	D/G	Sep	proposal, local outputs
12	Summary-native grounding	S	D/G	Oct	live (OOTA)

The table already shows the result of the survey: the approaches that a GM promoted, or that measured well, are those in which the precision decisions moved out of the model, and the two that read structured summaries (11 and 12) are the only ones in which the model is never asked to recover an order or an identity.

5 The approaches

5.1 Family P/M: extract from prose, synthesise with a model

1. Per-document scripts (March 2026). Four scripts—one per grounding document—each split the bible into chapter-aligned chunks of up to about 60,000 characters, ran a model extraction per chunk into their own extraction directory, and ran a second model call to synthesise the document from those extractions. A `--build-dossiers` mode of the planning script wrote per-NPC dossiers. The human role was review after the fact. The approach is simple and the documents read well, and three campaigns still use it. Its failures are of staleness and invention rather than style. An issue note from July records that “the extraction layer has Chapter 62. The synthesis layer did not consume it”: all four documents were stamped Chapter 62 and described Chapter 61. In another campaign the synthesis introduced a non-existent NPC that had to be removed by hand. Each script also re-reads the same input, which motivated approach 2.

2. Shared chapter extract (killed, April 2026). The consolidation replaced three narrow extractions with one rich extraction per chapter, in an eight-section schema, consumed by three synthesisers. At the extraction layer it looked better: the schema held across chapters of 2.7K–54K characters and one file per chapter was easier to review. At the synthesis layer every document

regressed. World state fell from 268 to 193 lines and lost its party and items sections; campaign state’s NPC table fell from 40 to 23 rows; planning fell from 513 to 109 lines, with 17 NPC dossiers collapsing to 5 and threat-tracker values “*guessed* (‘6–8 estimated’)” instead of read. Outputs were well under the token cap; this was depth loss, not truncation. The kill note’s diagnosis is general: “Consolidating N narrow LLM extracts into one broad extract is not free. The narrowness *is* the depth.” It also named the principle that recurs throughout this survey: “the rough extraction pass is the ceiling, not the floor.”

3. Hand curation. Underneath every generated approach sits human repair: merging duplicate dossiers (one pass merged 54 files into 25), sidecar files of new notes, candidate files for the party document so that a regeneration would not overwrite curated text, and hand fold-ins that carry a document forward from the latest session summaries without regenerating it. A commit message records the cost: dossiers “are generated and *then hand-curated*... Regenerating produces the machine output, not the curated state.” Hand fold-in is the current state of two campaigns’ documents, and it coexists uneasily with a stated rule that generated documents are not hand-edited.

5.2 Family F: atomic facts

4. Local atomic-fact ensemble (May–September 2026). To extract once and share the result, extraction moved to local hardware: five extraction *lenses* (prompts emphasising different material) times three samples, served from an 80B-parameter open model across two DGX Spark machines, emit atomic facts with a type and a subject. A deterministic merge groups facts by (type, normalised subject), or optionally by embedding similarity. A bundling step writes per-entity dossiers, with a listing mode documented as “the human checkpoint on what will be aggregated”. GM skills review alias and type merges. One model call per document then synthesises the grounding documents from selected dossiers. The first full run (1 June) produced 22,240 facts over 56 chapters and 703 dossiers; a parallel set of documents was written but, per its own commit, they were “not drop-in replacements... campaign_state is a derivation-of-derivation,” and that set was later deleted. One smaller campaign promoted ensemble drafts after GM review; the largest campaign never did.

A July investigation of the stale-chapter report found that the approach’s deepest problem was not a bug. “Atomization strips the syntax that carries attribution.” A sentence that identifies a dead antagonist is unambiguous in its paragraph and unrecoverable as a fragment; in the case that drove the investigation, three facts were extracted from one quote, two correct and one inverting the roles, and because bundling is by subject the inverted fact became the *only* fact in the antagonist’s dossier. Synthesis then amplified it to “Current status: Alive”. The obvious filter does not exist: 98% of facts (15,220 of 15,465) were produced by a single sample, because the lenses extract different material by design. Embedding similarity does not help either, since two statements that disagree about who died are far apart by construction (cosine 0.85 between the correct and inverted fact, below the merge threshold), and a contradiction detector found 15 candidate groups, 14 of them benign.

5. Repairs within the ensemble (July 2026). The investigation’s through-line was: “The pipeline computes a signal, serialises it, and then ignores it—forcing a later stage to guess what it was already told.” Four instances were fixed in a day. A dossier frequency floor ranked by fact count, which “deleted the newest chapter” because new entities have the fewest facts; a recency window restored all seven entities of the latest chapter. The embedding merge was unreachable from the interface, so “62 chapters silently merged with the weaker algorithm”. Quote positions

were reduced to a boolean, which lost within-chapter order; stamping offsets made the inverted fact visibly 800 characters after the scene that refutes it. Entities attested only through other entities’ facts were invisible; a coverage report surfaced 30, including one with zero facts of its own and 34 mentions that had no dossier at all. A chapter–summary mapping and a per-chapter narrative pass were added, both gated on explicit GM approval. These repairs made errors *legible*; they did not make the facts correct.

6. Two-stage extraction (proposal, September 2026). Replaying the shipped merge over a fresh corpus of 13,731 facts measured what the prose comparator costs: 2,384 pairs of rows whose source spans overlap but were kept separate (821 with byte-identical quotes), and 43 merged rows whose members sit on disjoint spans—in one, three of a chapter’s own date headings fused into a row asserting one date. The proposal splits extraction into a finder stage that produces spans and a claim stage that cites them, so the span becomes the merge key. An open-source grounded-extraction library was evaluated as the finder and rejected: on facts it was not stable across chapter size, falling to 6.2% source coverage on a 14K-character chapter against 75.6% for the ensemble union, as extractions degenerated into bare names.

7. Entity-level typing and ruling replay (October 2026). Typing each fact independently splits an entity across types (an NPC in one chapter, a monster in the next); the GM had resolved 109 such splits by hand in one campaign. A companion study [3] found that re-typing facts with a non-generative decision model “creates roughly as many splits as it repairs.” The replacement types once per entity through a deterministic chain: GM filing conventions, then prior GM merge rulings, then the registry type, then a 60% majority vote, then a GM queue. The model-free vote matched the GM’s filing 32 times in 33 where it decided. The proposal’s sharper conclusion was that the pipeline “asks the LLM for precision decisions and never reads the GM’s rulings back”: rulings were recorded and then ignored on the next run.

8. Model work layered on ensemble dossiers (July–August 2026). Several local jobs, run by a locally served model through an agent harness and kept outside version control, worked on the ensemble’s dossiers: appending matching facts to each dossier; answering each dossier’s uncertainty bullets from the cited chapters (327 dossiers: 158 resolved, 140 partial, 11 unresolved, 18 malformed outputs; four dossiers found contradicted by the chapters); per-entity chapter scans and an inferred calendar timeline; and a cross-campaign “scholar persona” wiki with a sampling fact-checker. None was promoted. Their main yield was evidence that unreviewed model digests drift on names and that dossiers built from atomic facts carry contradictions that a later pass can find but not settle.

9. State stores and projections (July–August 2026). This approach changed the output form rather than the extraction. Durable stores—a deduplicated per-chapter event spine, a registry of GM-ratified plot threads with model-proposed candidates, and a GM-approved chapter–session map—feed per-section renders, and a document is the assembly of its sections. A build without a model backend makes “zero model calls; a re-run with unchanged stores rebuilds nothing.” Planning is layered by certainty, from ratified threads down to explicitly labelled non-canon speculation. The design is sound; its inputs were not. The event spine was built from ensemble facts, and near-duplicate phrasings of one event survived its 60-character deduplication key. In the one campaign that used it, four threads had been ratified and 148 proposals were pending when work stopped, and that campaign’s live world state returned to hand fold-in.

5.3 Family R: retrieval instead of documents

10. Knowledge graphs and semantic search (July–September 2026). A knowledge-graph library [4] was used to ingest grounding documents and dossiers for one campaign and session summaries for another, with a locally served model building the graph. The first attempt never completed in about six runs; a later retest completed in 27 minutes and answered five evaluation queries correctly, but the graph was never adopted, because the structure was written by a model with no review. The September investigation explains why that matters. The library stored 2,468,434 characters of generated text for 2,093,539 characters of source (1.18×), including paraphrase drift, facts imported from the published adventure that never happened at the table (327 edges about a castle the party has never found), and a false identity merge of a villain with his master. A stripped prompt did not stop it: with structured output, “the schema is the stronger instruction.” Asked for the most recent session of an eleven-session campaign, the graph returned Chapter 8, a vector search tool returned session 6 and a semantic memory store [5] returned session 8. “Every semantic tool failed the time question.” The verbatim tools never fabricated; their failures were in ranking. A semantic search skill over the summaries remains live as a *lookup* aid alongside the documents, with filename ordering and an explicit absence check added because “vector search structurally cannot” supply them.

5.4 Family S: authored structure

11. Event-ordinal graph (September 2026). The September investigation’s conclusion was that “the sequence of events is already authored”: filenames give session order, each summary’s **Scenes** section gives scene order, and its bullets give beat order. “Time is the event ordinal, not the calendar.” A graph keyed on that ordinal, with state snapshots from the summaries’ entity sections and identity taken from the GM-reviewed alias file (“the one precision decision”), was built in about a second for twelve sessions (72 events, 936 beats, 253 snapshots), answered every time question the semantic tools failed, and was reproduced byte-identically by three independent extractions. Retargeted to a 53-session campaign it ran in 0.8 s (356 events, 5,255 beats). Identity resolution through the alias file raised one NPC’s mentions from 4 to 15 and, unlike the model-built graph, did not merge the villain with his master. The investigation also observed that the alias file mixed surface-form equivalences with secret identities known only to the GM, and that an identity revelation is itself an event. Deterministic state documents were generated from the graph by exact search and copy, with no model. They were never promoted, and the proposal to replace the per-document campaign-state script was not built.

12. Summary-native grounding (October 2026). The approach now in use for the largest campaign reached the same design independently, starting from a failed attempt to run the ensemble over summaries instead of prose (a four-lens extraction projected at 7–10 hours on the available hardware). It reads a directory of structured summaries and nothing else. Validation runs first, collects every problem across the corpus, and refuses to proceed on any error in range: filename number, title number and scene identifiers must agree, and a file without a **Scenes** section is never treated as one undifferentiated block. A deterministic parse then produces a chronology of every scene, the memorable-moments sections preserved character for character, and one dossier per entity with the source file and scene of every observation. Headings are grouped only when identical within a category or listed as exact same-type aliases in the registry. Likely duplicates (similarity ≥ 0.88 , or a name with and without a parenthetical qualifier) are *listed, never merged*: “No merge decision MAY be made by a model or a similarity score”—the GM corrects the summary, or records the pair

as distinct in a rulings file the pipeline only reads. Every deterministic artifact is byte-identical across runs given identical inputs.

Only then does a model write: world state, campaign state, party and planning drafts, from dossiers selected by recency and recurrence and from the full chronology. Tracking and planning files are given to campaign-state synthesis as *questions to audit*, not as evidence, and any claim the summaries do not support is labelled rather than narrated. A draft feeds the next synthesis only when the GM names it as reviewed; drafts never overwrite live documents; exact prompts, input digests and selection decisions are retained. On the full corpus the parse covered 67 of 67 summaries, 409 scenes and 868 dossiers; the drafts extended the live documents from Chapter 63 to Chapter 70; and the campaign-state audit marked 51 tracking claims as not found in the summaries—planned or module-default events that never happened at the table. The GM judged the results “much better than the other ones we have built” and promoted all four.

5.5 Cross-cutting: the entity registry

Every approach from 4 onward depends on the entity registry for identity. It was built from module inventories through a model pass, a deterministic merge and GM rulings, imported by tooling rather than edited by hand (954 entities in one campaign, “no manual YAML surgery”), and later cut from 990 to 922 entries in another by removing misspellings that had been recorded as aliases. The rule that emerged is that the registry records names, types and aliases, never state or history.

6 Findings

F1. The extraction pass is the ceiling, not the floor. An impressive first-pass extraction predicts nothing about the synthesis built on it (approach 2), and a synthesis cannot recover detail that the extraction never wrote down at the needed granularity.

F2. Atomisation destroys attribution. Decontextualised facts are what make deduplication and counting possible, and they are also why “who did what to whom” cannot be recovered downstream by clustering, voting or embedding (approach 4).

F3. Agreement is not available as a filter. Ensembles are usually justified by agreement. With lenses designed to extract different material, 98% of facts had a single supporting sample, and embeddings measure paraphrase, not contradiction.

F4. Computed signals get discarded, and the system then guesses. Chapter ranges, merge method, quote positions and mention counts were all computed and then ignored by the next stage (approach 5). The repair in each case was to carry the signal forward, not to add a model.

F5. Silent success is the dominant failure mode. A library ingested a file path as text and reported success; a merge fell back to a weaker algorithm with a message that read as configuration; a reasoning model spent its whole completion budget thinking and returned no content; three tools reported success over work they had not done. Instruments need positive controls.

F6. Time is an ordinal, and semantic retrieval cannot find it. Every semantic tool failed “most recent session”. Calendar dates are sparse and inconsistently formatted; the authored scene order is always present (approaches 10, 11).

F7. Generative structure invents the record. Graph-building and summarising layers import module canon as history, drift in paraphrase and merge identities, and store the result as fact (approach 10). Verbatim tools fail by omission, which a reviewer can detect; generative layers fail by commission, which a reviewer must catch.

F8. Identity is the precision decision. Which names denote the same entity is the one decision that cannot be read off the text, so it belongs to the GM, recorded once and read back every run (approaches 7, 11, 12). Surface-form equivalence and secret identity are different decisions and need different scopes.

F9. Plans are not evidence. Tracking files, planning notes and module text describe what was expected. Treating them as audit questions exposed 51 claims that never happened (approach 12).

F10. Do not undo a rendering step. The narrated bible is rendered from the structured summaries. Re-extracting structure from it is a lossy round trip—“structure $\rightarrow$ render to prose $\rightarrow$ re-extract structure from the prose”—and the dimension it loses (order, attribution) is the one the extractors then spend tokens guessing. The cheapest, fastest and most accurate approach read the structure that already existed.

7 Discussion

Where the model belongs. Across the twelve approaches the model’s role shrank from extractor, structurizer and renderer (approaches 1–2) to renderer only (approach 12), and each step of that shrinkage was forced by a measured failure, not by principle. The final division is: authored structure supplies order; the registry and recorded GM rulings supply identity; deterministic code supplies grouping, selection and provenance; the model supplies prose. The human checkpoint did not move—it was always the GM—but it moved *earlier*, from reviewing generated documents to reviewing the summaries and rulings the documents are computed from, where a correction is made once and persists.

Rulings must be inputs. Several approaches asked the GM for a decision, recorded it, and asked again on the next run. A ruling store that every stage reads—conventions, merges, aliases, not-a-duplicate pairs—is the mechanism that keeps a checkpoint from becoming a recurring chore that an exhausted reviewer eventually rubber-stamps.

What remains open. Summary-native grounding moves the precision problem upstream; it does not eliminate it. The summaries are themselves produced with model assistance and reviewed by the GM, and their quality bounds everything downstream; chapters without recordings need summaries written from the narrated prose, reintroducing a rendering step. The store-and-projection design (approach 9) is the natural way to make summary-native outputs incremental and section-addressable, but it was built over ensemble facts and has not been retargeted. Three campaigns still run approach 1, and migrating them requires their summaries to conform to the structured schema. Finally, the history itself was at risk: the prototype that led to approach 12 was built in a temporary worktree whose code was never committed, and the event-ordinal outputs and several local experiments exist only on one machine. A practitioner running many exploratory approaches should commit the evidence even when the code is disposable.

Cost and local hardware. The local ensemble was attractive because extraction on owned hardware is nearly free per run. It was not free in wall-clock time (hours per corpus, 7–10 hours projected for one summary run) or in review time (tens of thousands of facts no one reads). The deterministic approaches run in seconds. The decisive saving was not cheaper extraction but not extracting.

8 Limitations

This is a retrospective survey of one practitioner’s work, reconstructed by an AI agent from repositories and documents written during the work; it inherits their framing. The approaches were not evaluated on a common benchmark: measurements were taken on different campaigns, corpus versions and models, and comparisons across approaches are qualitative. The GM’s judgement of the final documents is a single expert’s assessment, not a blind comparison. Status labels for campaigns not examined in depth (in particular which method produced one campaign’s latest refresh) are uncertain. Local-only work was assessed from its outputs and scripts, without commit history. Summary-native grounding has been promoted for one campaign; its behaviour on campaigns with less regular summaries is untested.

9 Conclusion

Twelve attempts to build grounding documents for long-running campaigns, spread over seven months and six campaigns, converge on a simple conclusion: the record a GM needs is already written, in the structured session summaries the GM reviews after every session, and the job of tooling is to read that structure faithfully, keep identity decisions with the GM, and let a model do only what it is good at, which is writing. The approaches that asked a model to recover order, attribution or identity from prose or fragments failed silently and compounded their errors; the approaches that read authored order and recorded rulings were faster, reproducible and better, by the only judge whose opinion counted.

Data and code availability

The toolkit, including the design documents, experiment write-ups and the summary-native pipeline, is public [1], as are the local-model experiments [2]. The campaign data is not released because it quotes play sessions verbatim and names real people.

Use of AI assistance

The approaches surveyed were largely implemented by AI coding agents (Claude, Anthropic, and others) under the author’s direction. This survey was compiled by an AI agent from the repositories’ history and documents, and drafted by it; the author made the decisions the paper records and takes full responsibility for the content.

References

- [1] K. Roussos. CampaignGenerator: tooling for tabletop campaign management with language models. <https://github.com/kostadis/CampaignGenerator>, 2026.

- [2] K. Roussos. dgx-fun: experiments with local LLM serving on NVIDIA DGX Spark. <https://github.com/kostadis/dgx-fun>, 2026.
- [3] K. Roussos. Decision models at the game table: where non-generative classifiers help and fail in LLM tooling for tabletop role-playing campaigns. Manuscript, October 2026.
- [4] Topoteretes. Cognee: memory for AI agents. <https://github.com/topoteretes/cognee>, 2026.
- [5] MemPalace contributors. MemPalace: a semantic memory palace for agents. Open-source software, 2026.
- [6] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W. Yih, T. Rocktäschel, S. Riedel, and D. Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*, 2020.
- [7] D. Edge, H. Trinh, N. Cheng, J. Bradley, A. Chao, A. Mody, S. Truitt, and J. Larson. From local to global: a graph RAG approach to query-focused summarization. arXiv:2404.16130, 2024.
- [8] C. Packer, S. Wooders, K. Lin, V. Fang, S. G. Patil, I. Stoica, and J. E. Gonzalez. MemGPT: towards LLMs as operating systems. arXiv:2310.08560, 2023.
- [9] J. S. Park, J. C. O’Brien, C. J. Cai, M. R. Morris, P. Liang, and M. S. Bernstein. Generative agents: interactive simulacra of human behavior. In *Proceedings of UIST*, 2023.
- [10] Y. Chang, K. Lo, T. Goyal, and M. Iyyer. BoookScore: a systematic exploration of book-length summarization in the era of LLMs. In *International Conference on Learning Representations*, 2024.
- [11] N. F. Liu, K. Lin, J. Hewitt, A. Paranjape, M. Bevilacqua, F. Petroni, and P. Liang. Lost in the middle: how language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173, 2024.
- [12] C. Callison-Burch, G. S. Tomar, L. J. Martin, D. Ippolito, S. Bailis, and D. Reitter. Dungeons and dragons as a dialog challenge for artificial intelligence. In *Proceedings of EMNLP*, 2022.
- [13] A. Zhu, K. Aggarwal, A. Feng, L. J. Martin, and C. Callison-Burch. FIREBALL: a dataset of dungeons and dragons actual-play with structured game state information. In *Proceedings of ACL*, 2023.
- [14] A. Zhu, L. J. Martin, A. Head, and C. Callison-Burch. CALYPSO: LLMs as dungeon masters’ assistants. In *Proceedings of AIIDE*, 2023.
- [15] Z. Ji, N. Lee, R. Frieske, T. Yu, D. Su, Y. Xu, E. Ishii, Y. J. Bang, A. Madotto, and P. Fung. Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12), 2023.